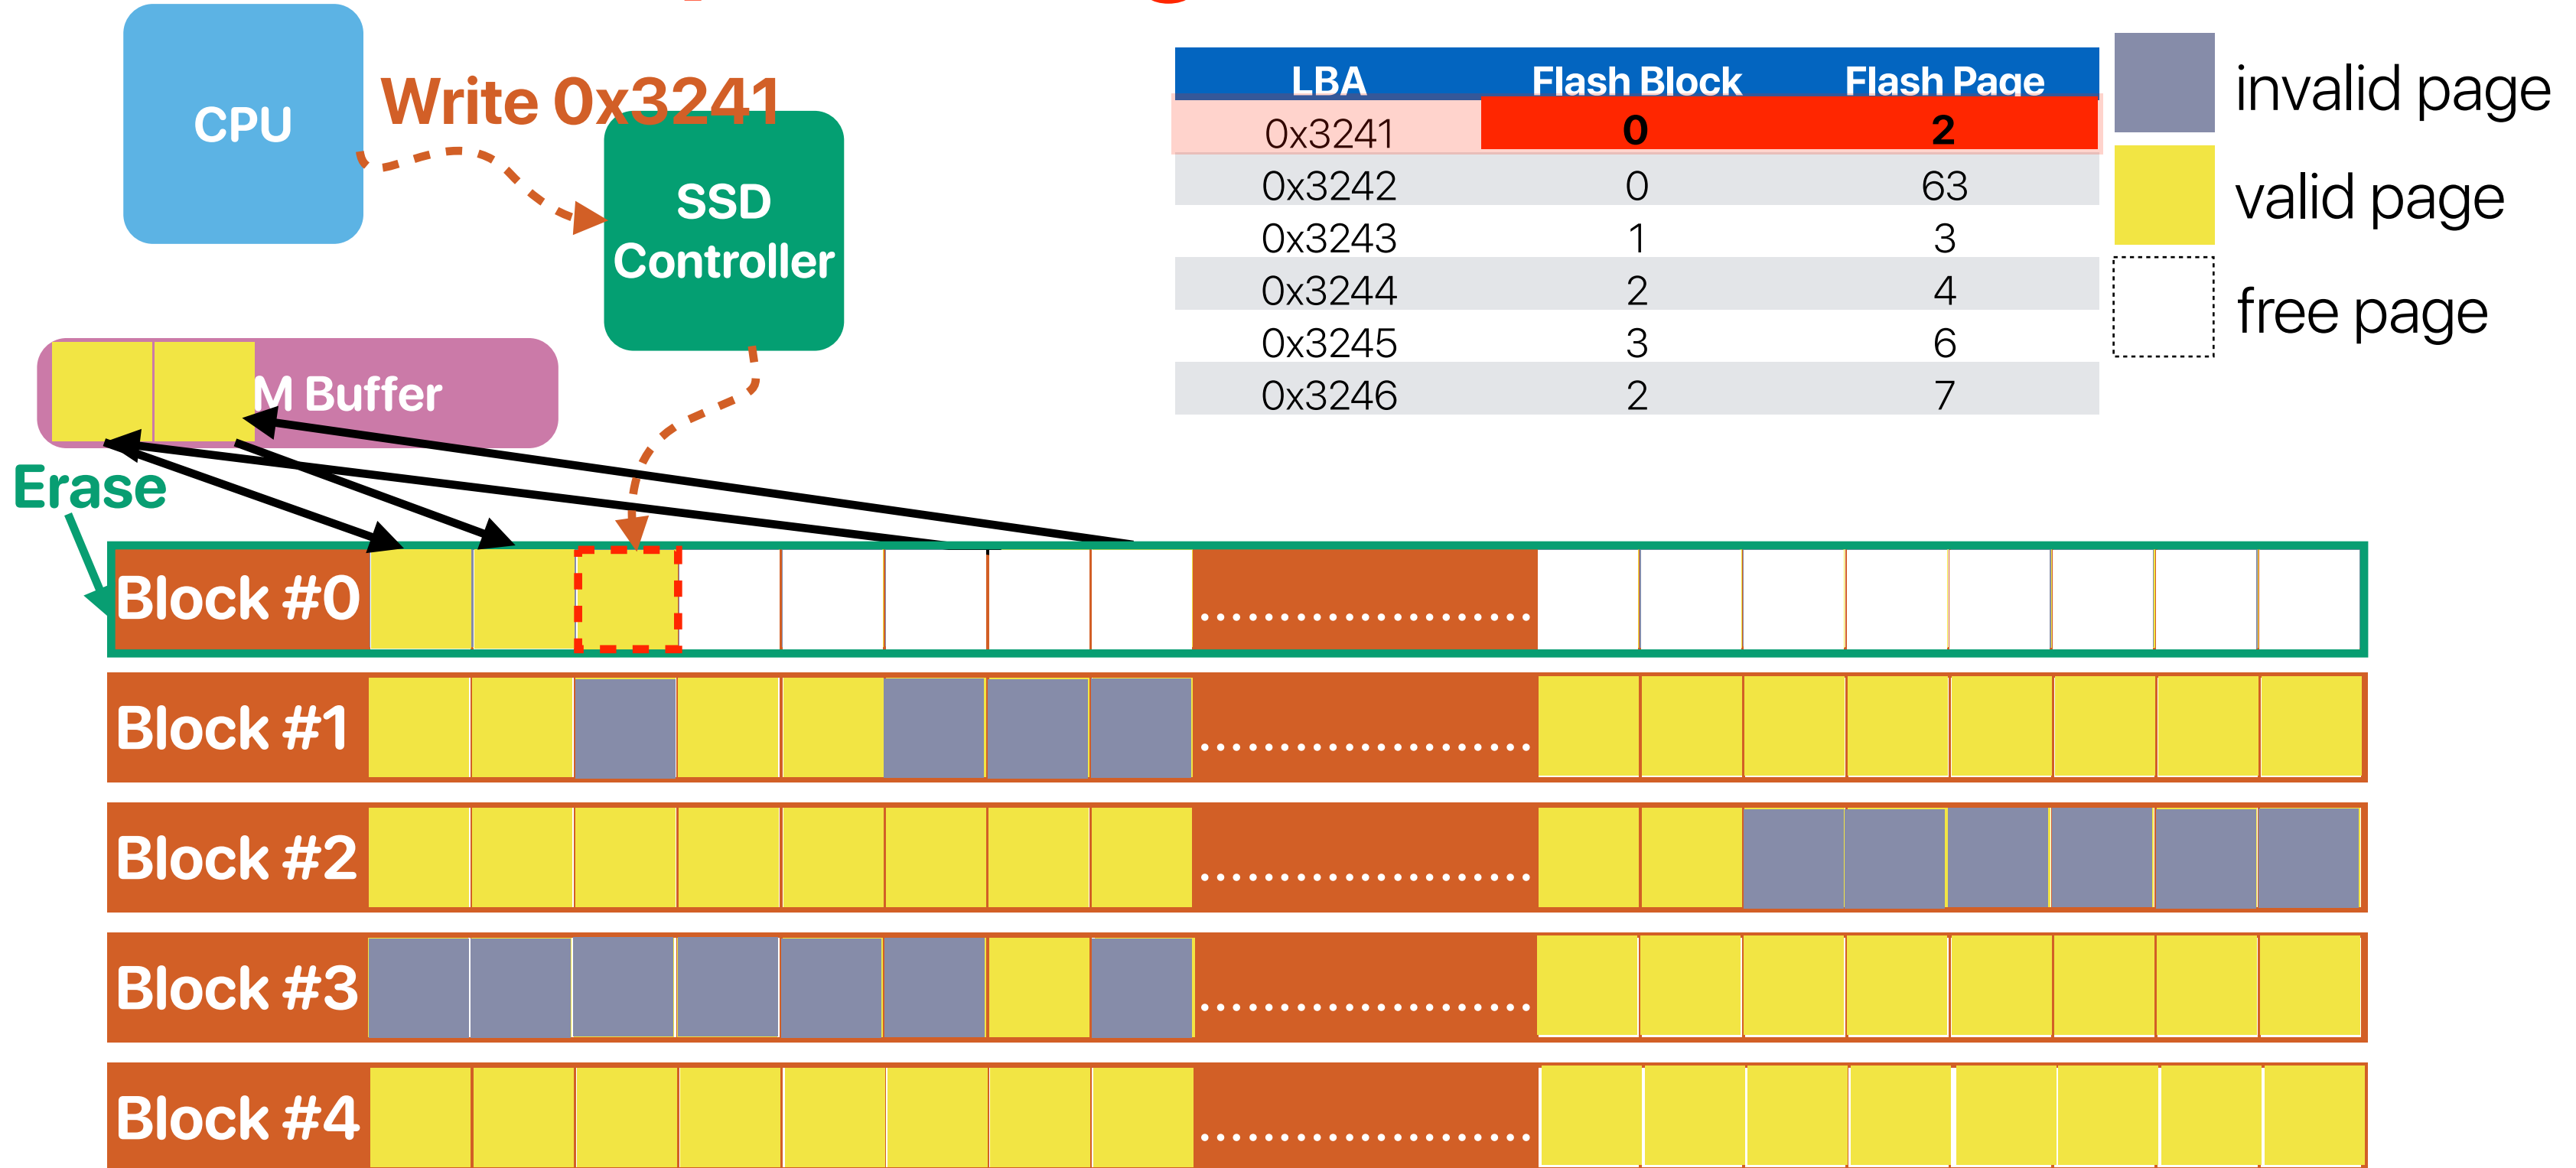


Streamline Data Exchanges & In-Storage Processing


Hung-Wei Tseng

Recap: Garbage Collection in FTL

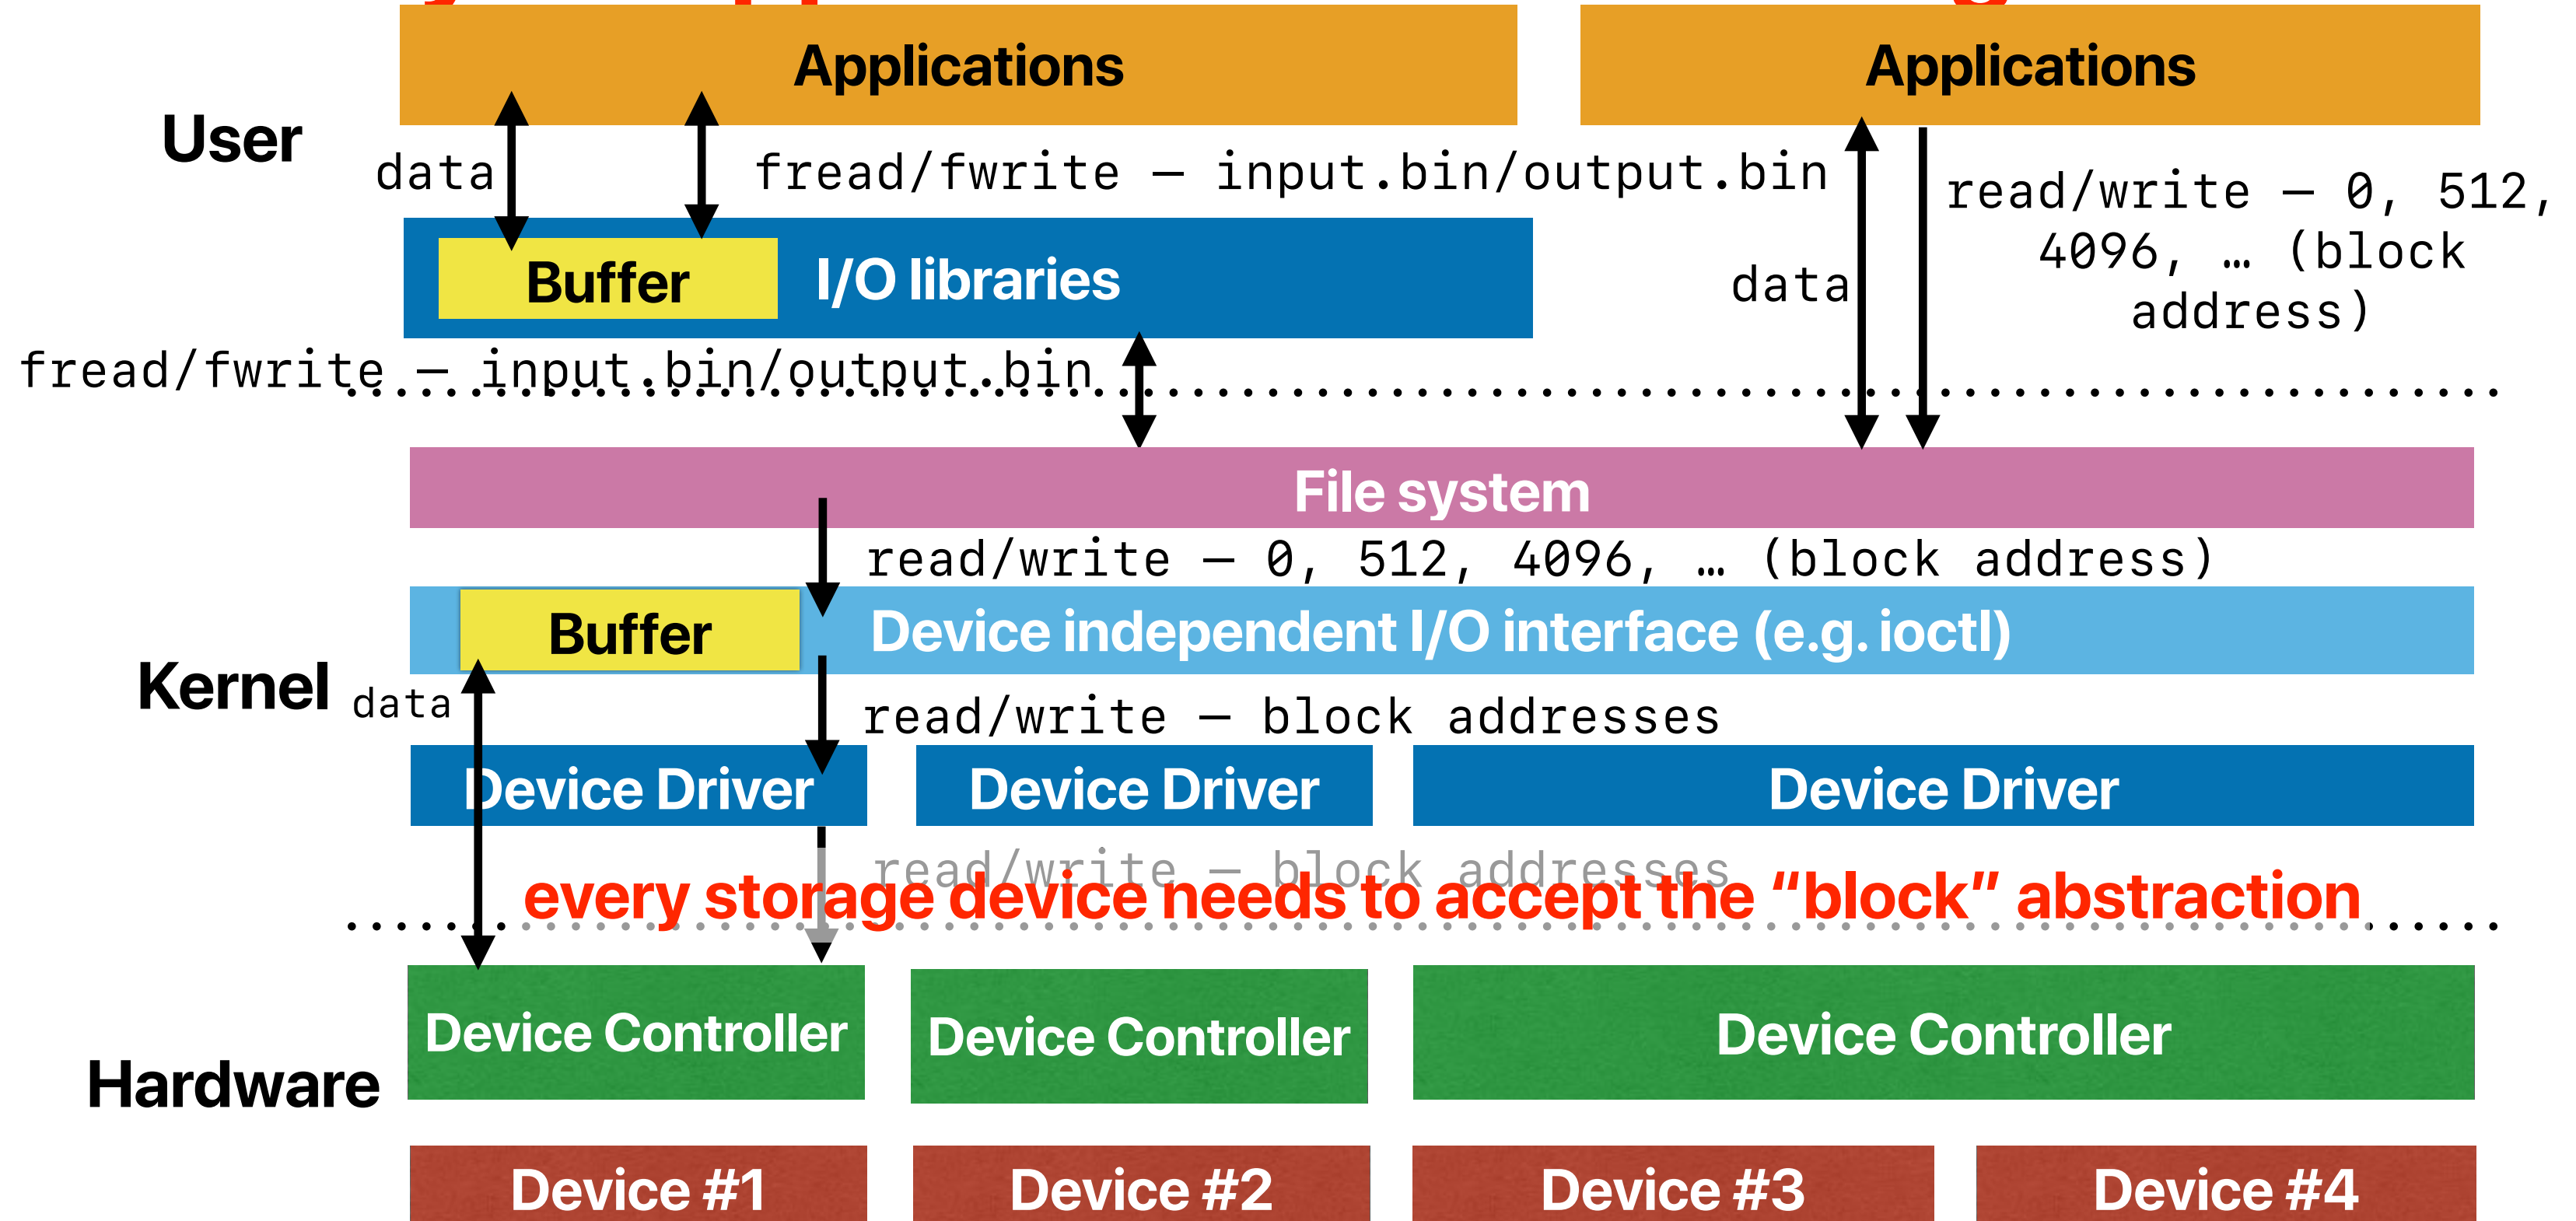


Recap: NVMe is the payload/data of PCIe

- Very similar to computer networks
- Use "memory addresses" as the identifier for routing

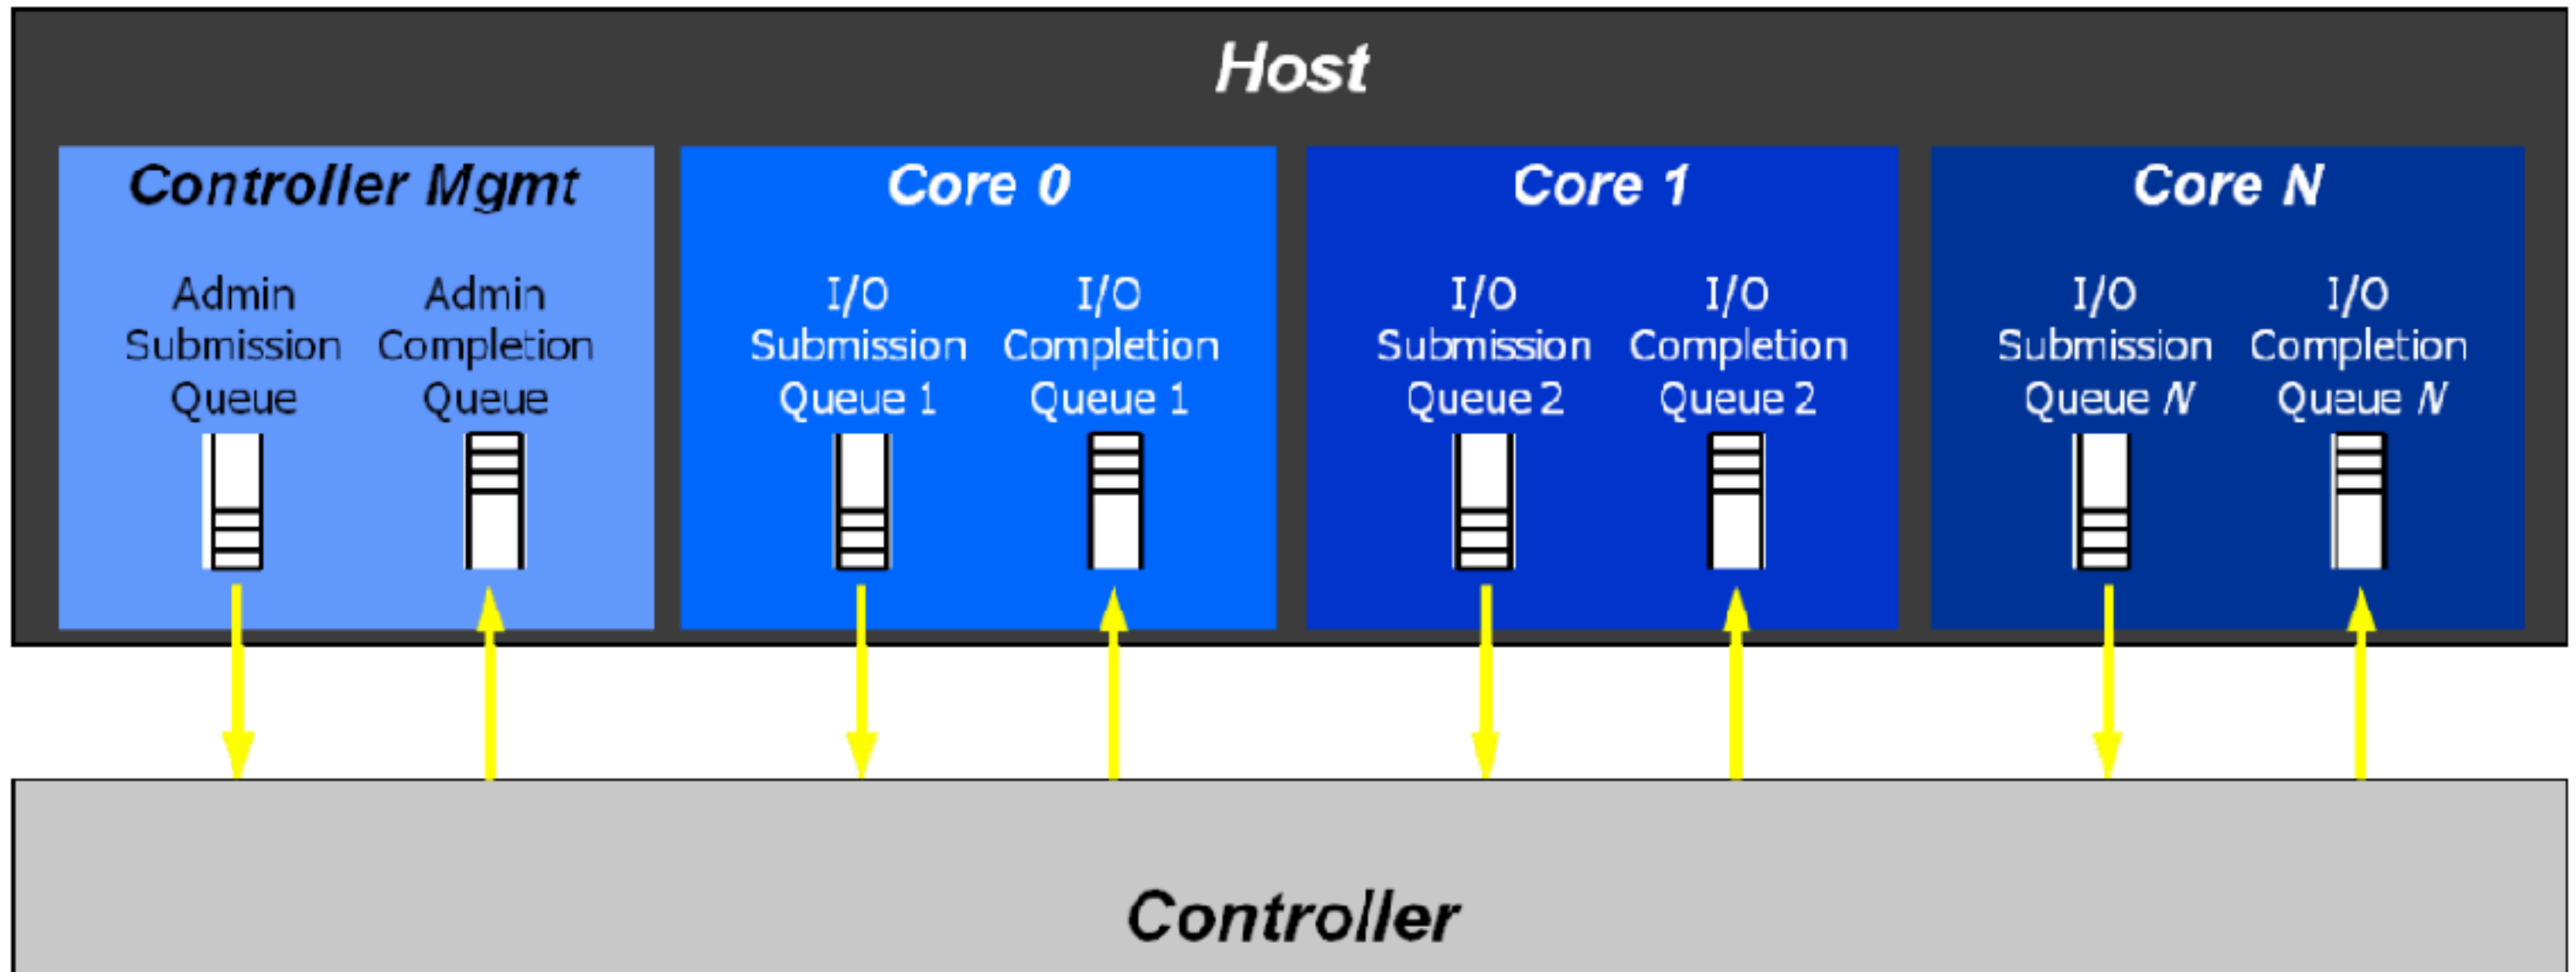


How your application reaches storage devices




Recap: NVMe Model

Figure 3: Queue Pair Example, 1:1 Mapping



https://nvmexpress.org/wp-content/uploads/2013/04/NVM_whitepaper.pdf

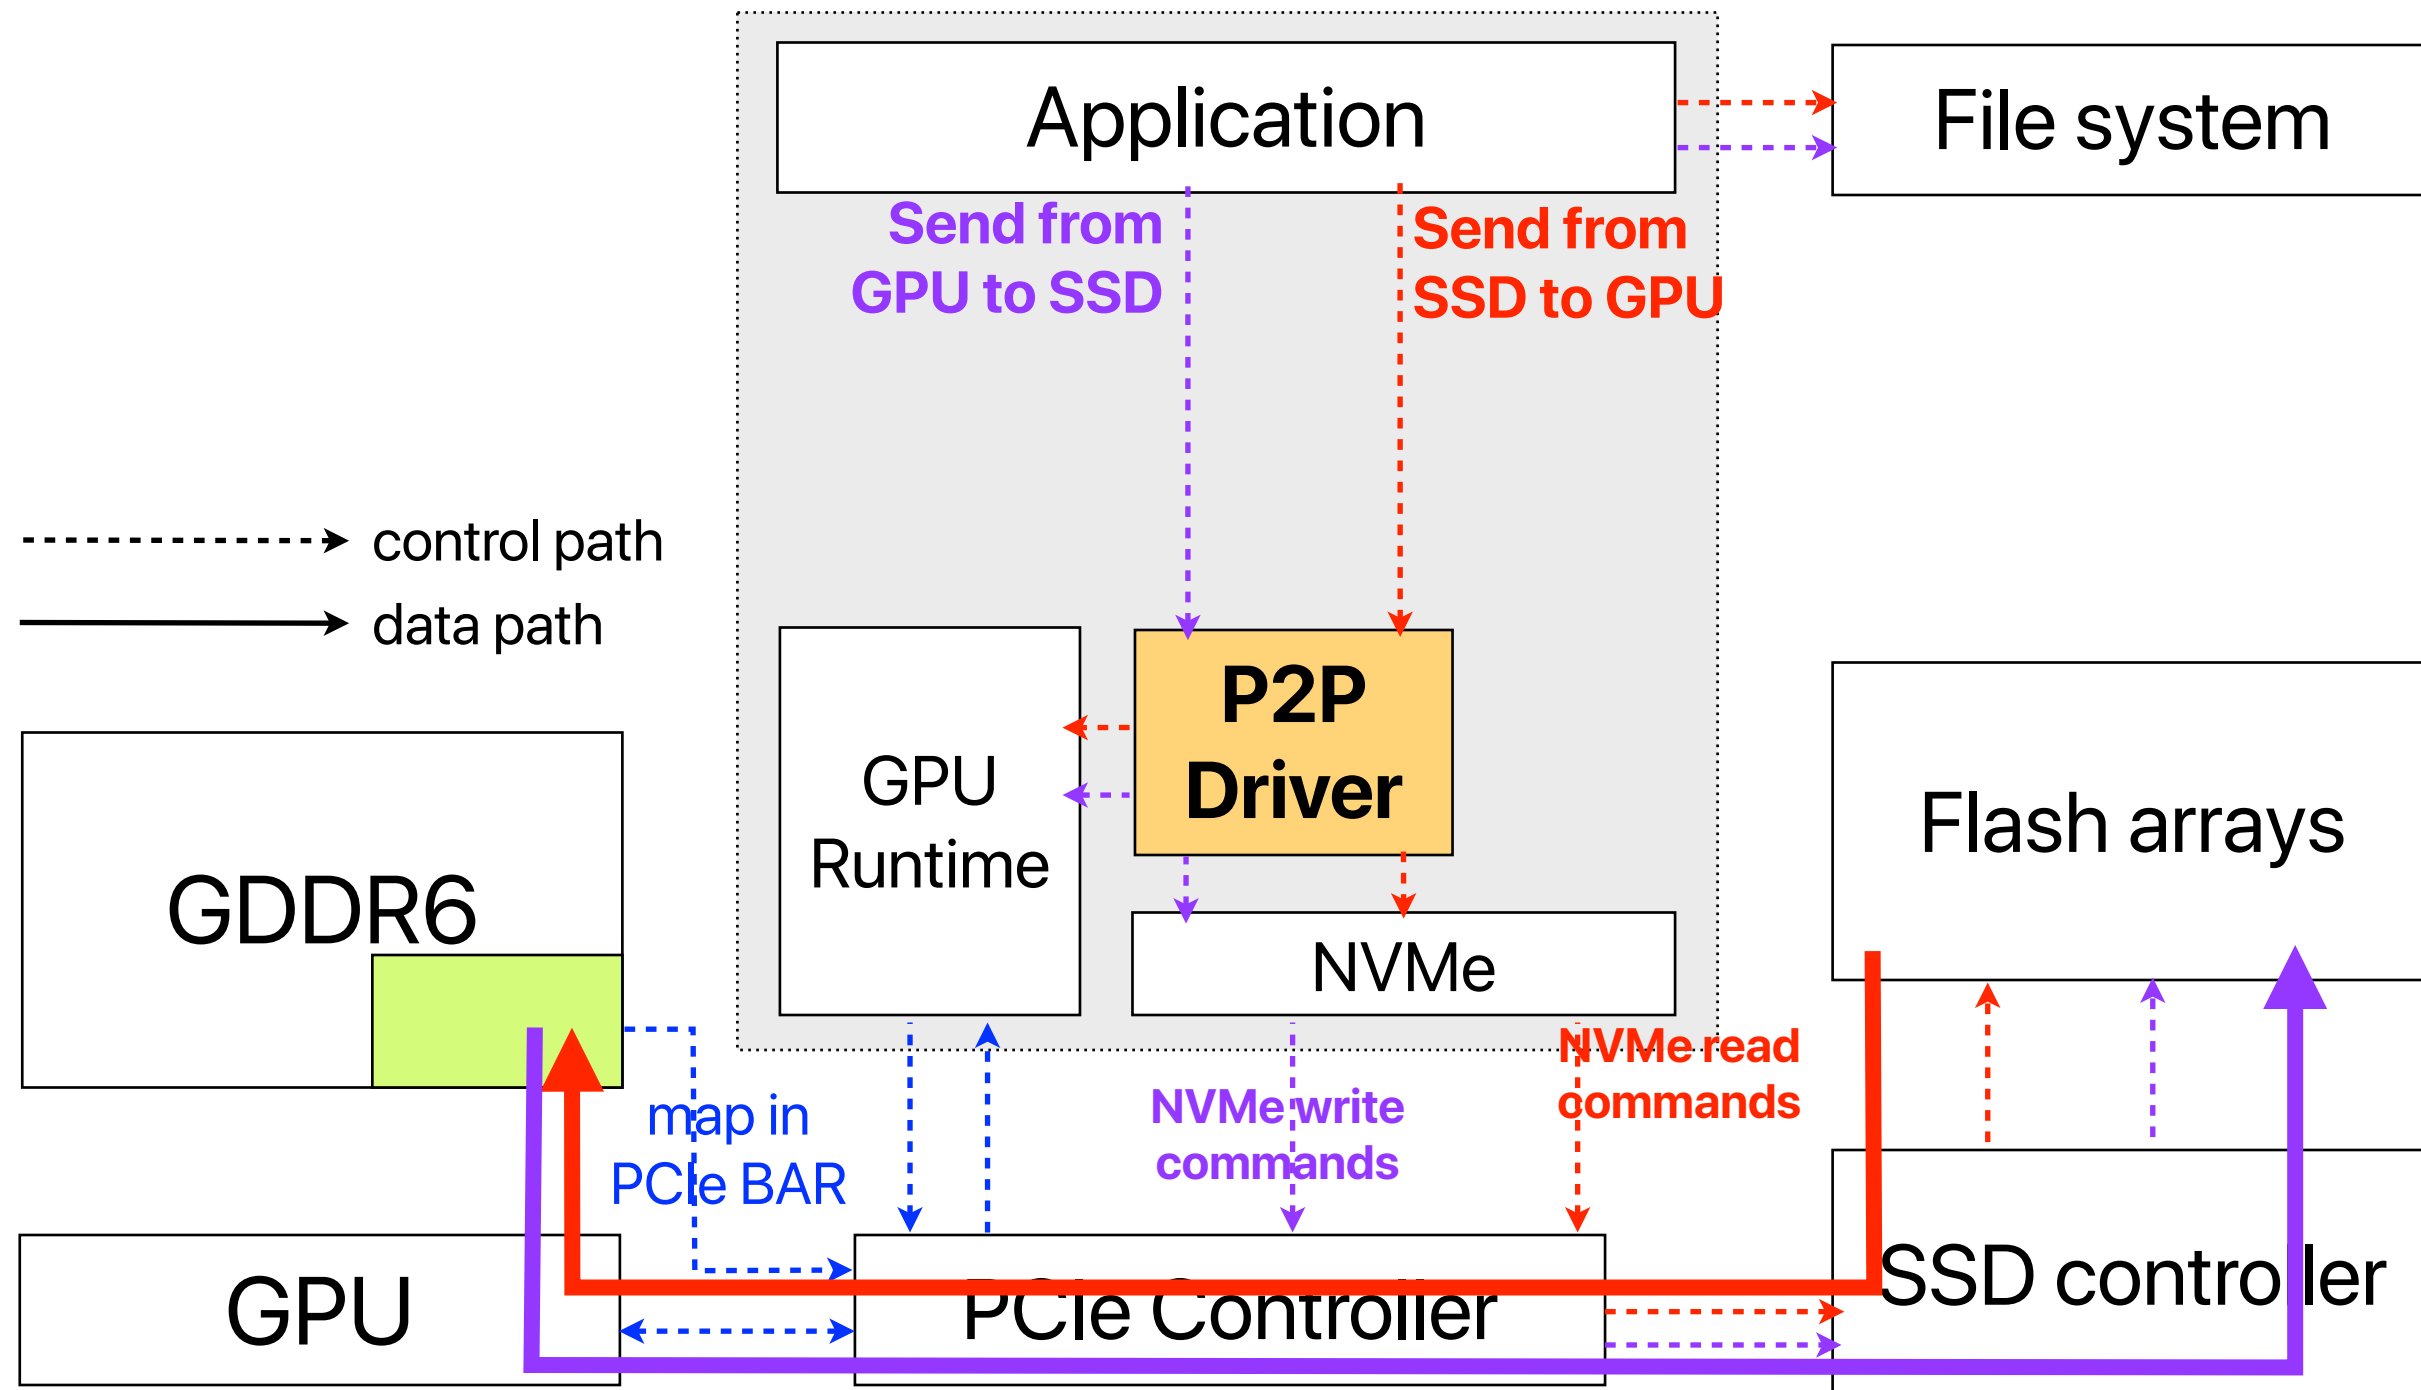
Recap: PCIe "Interconnect"



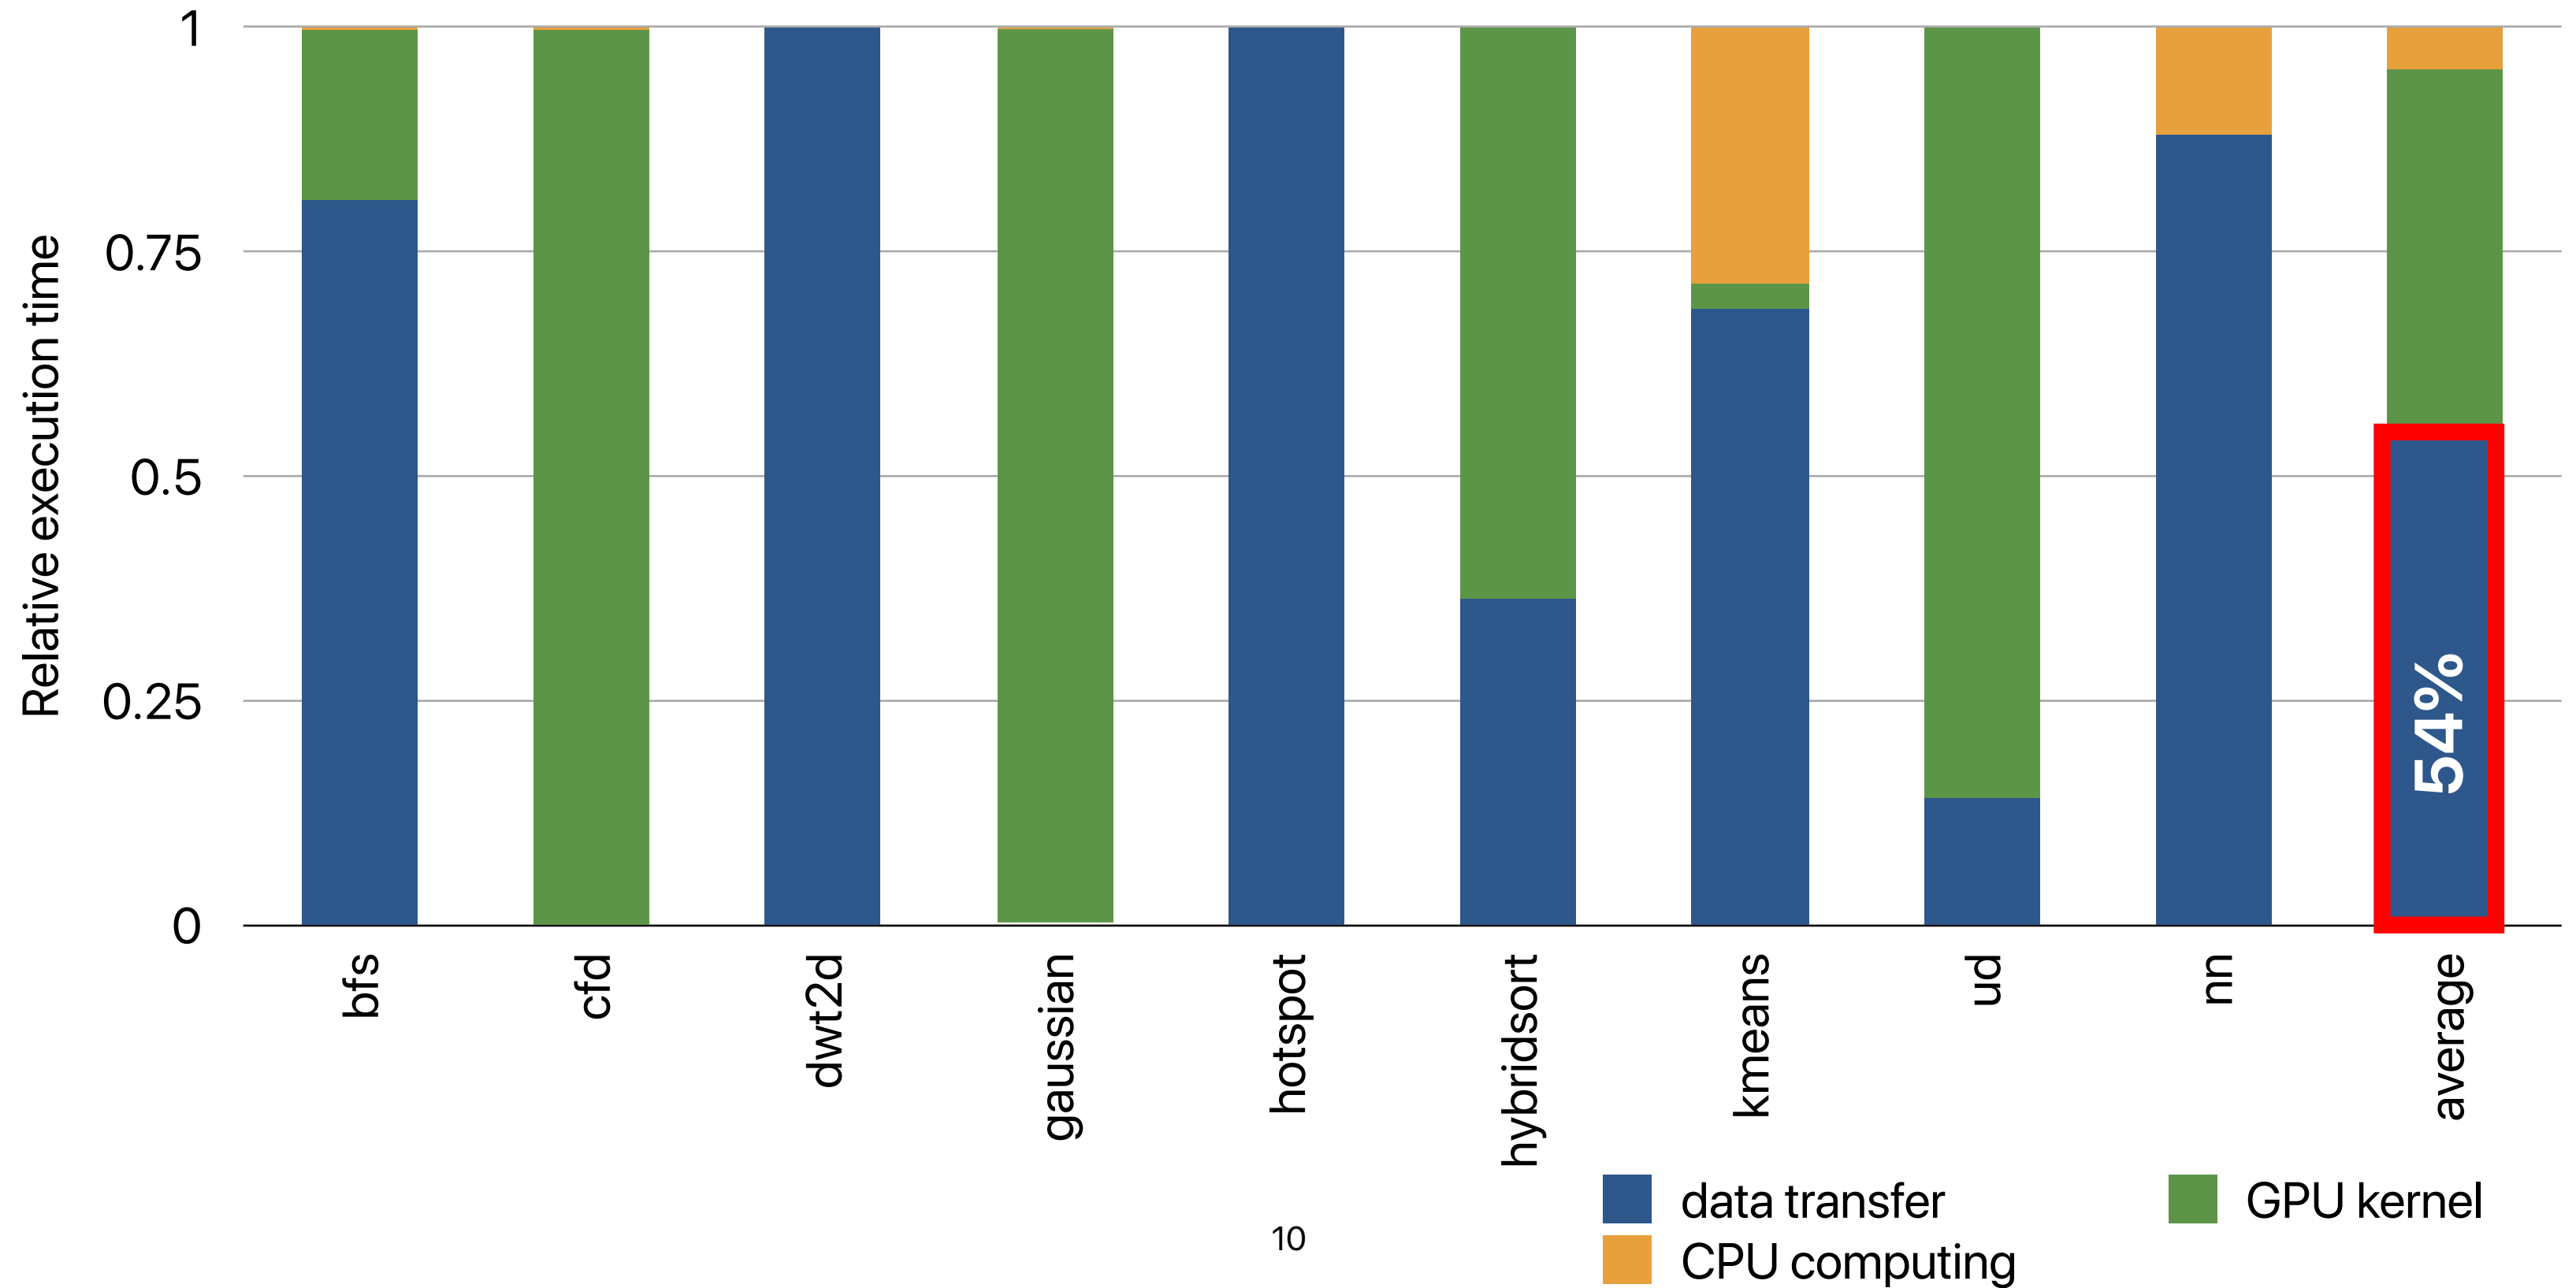
**How can you let SSDs to
communicate with GPUs directly &
why does SSD-GPU P2P make sense?**

How to implement GPU-SSD P2P?


P2P between GPU and SSD



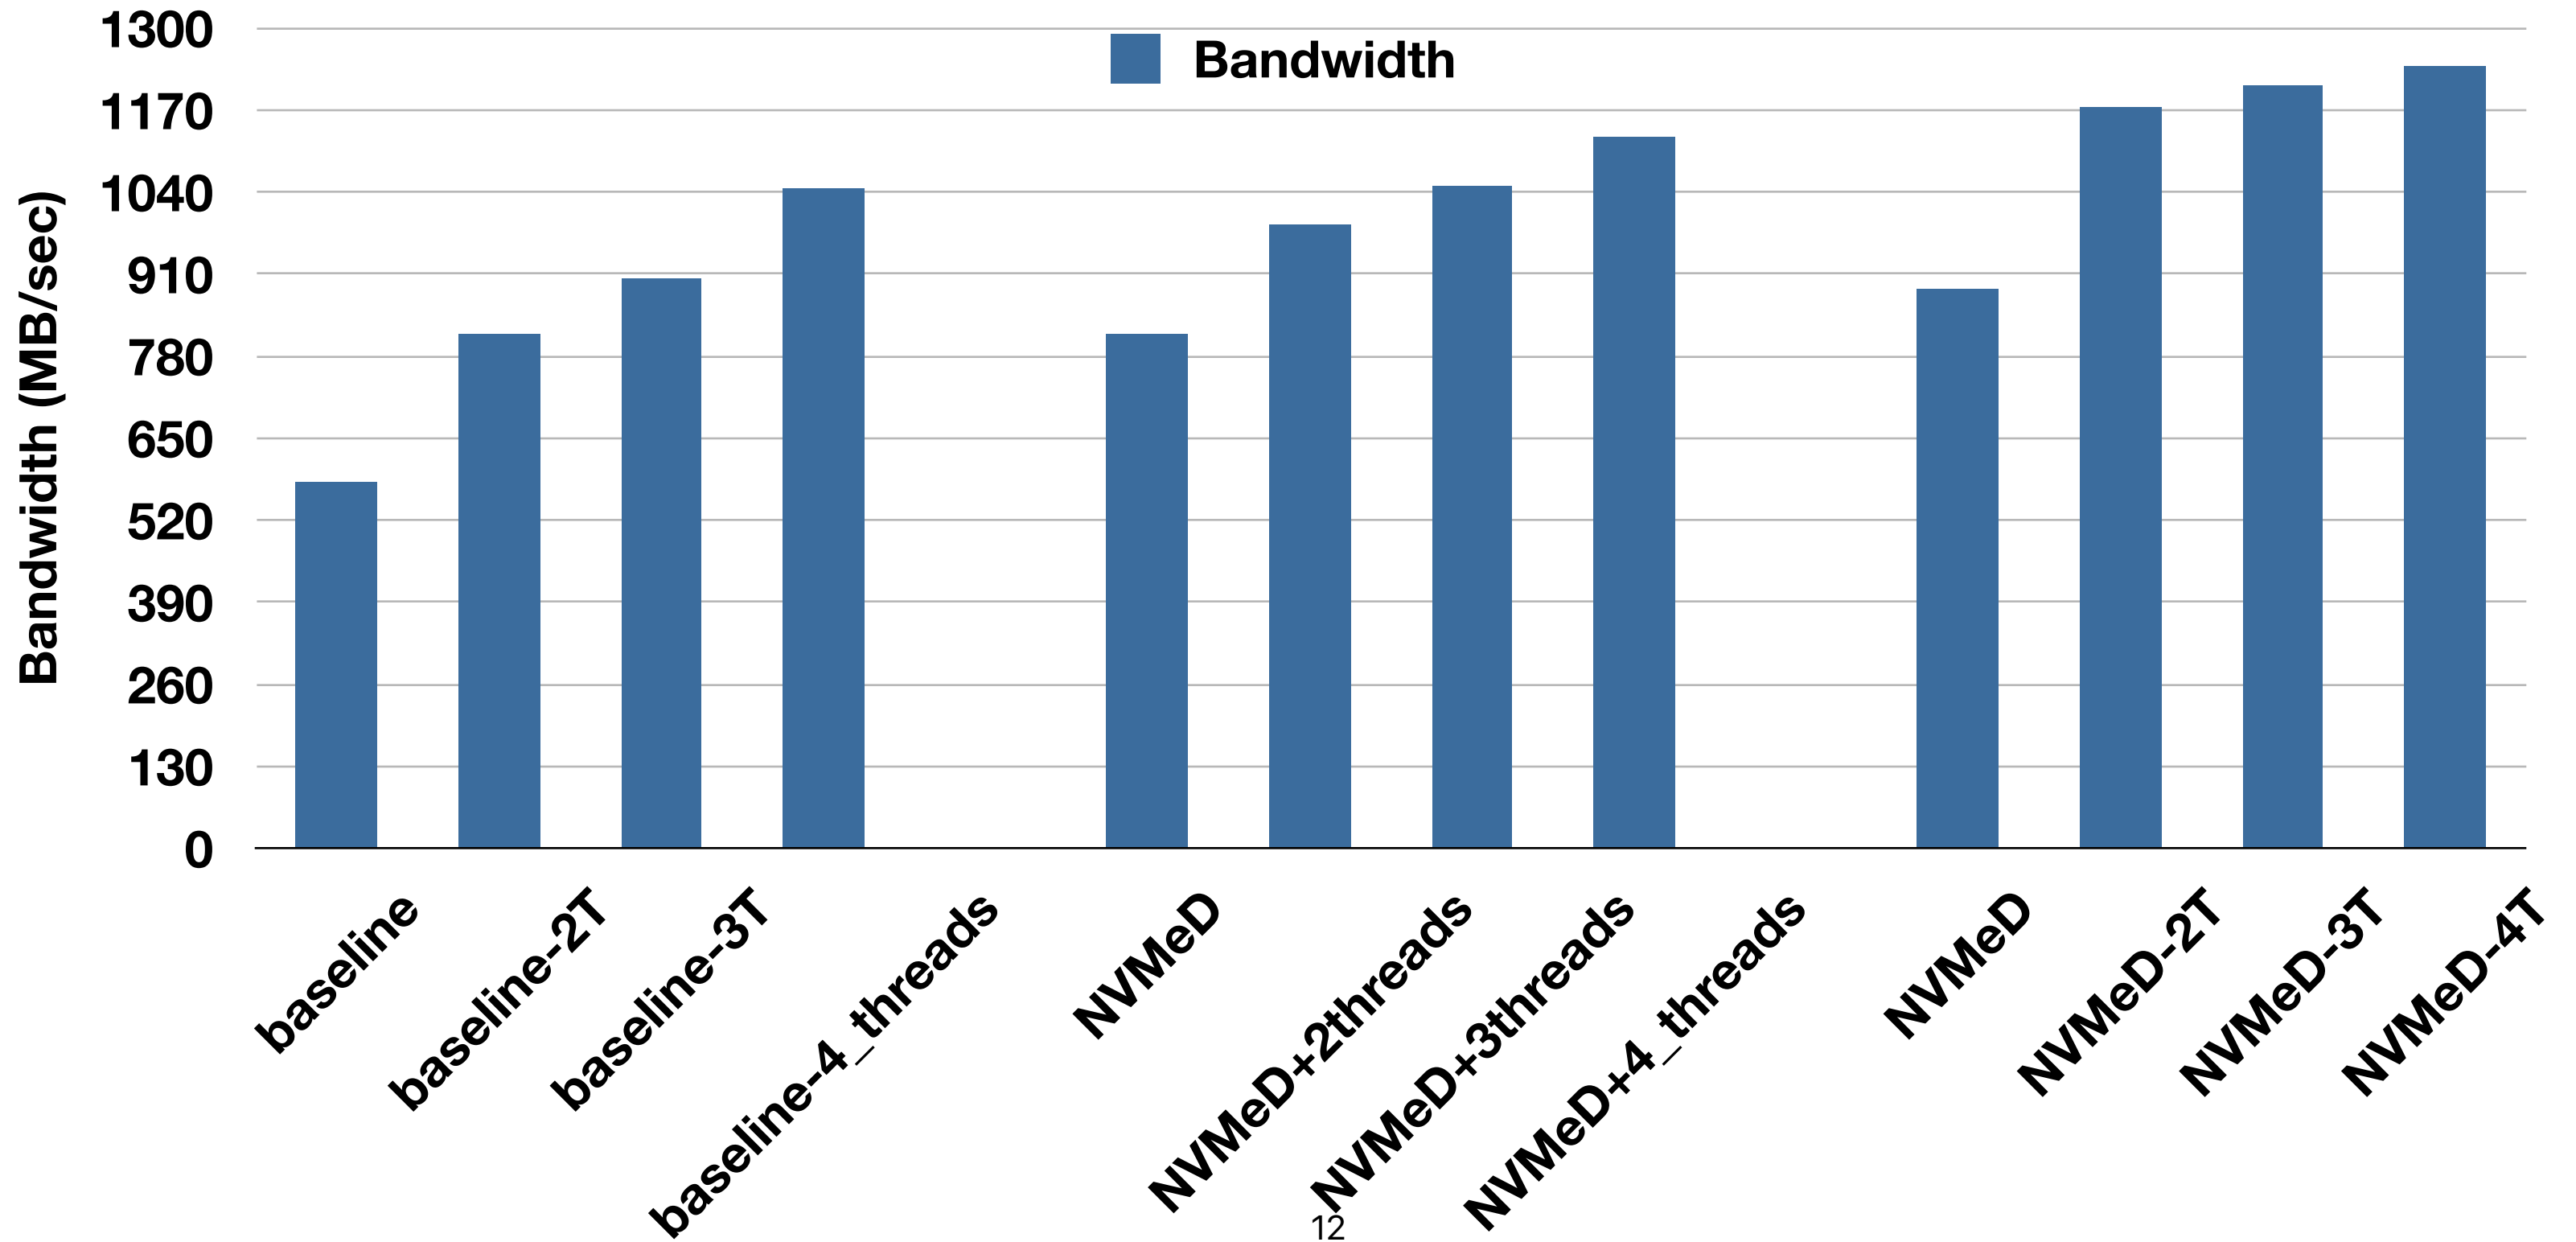
Cost of moving data



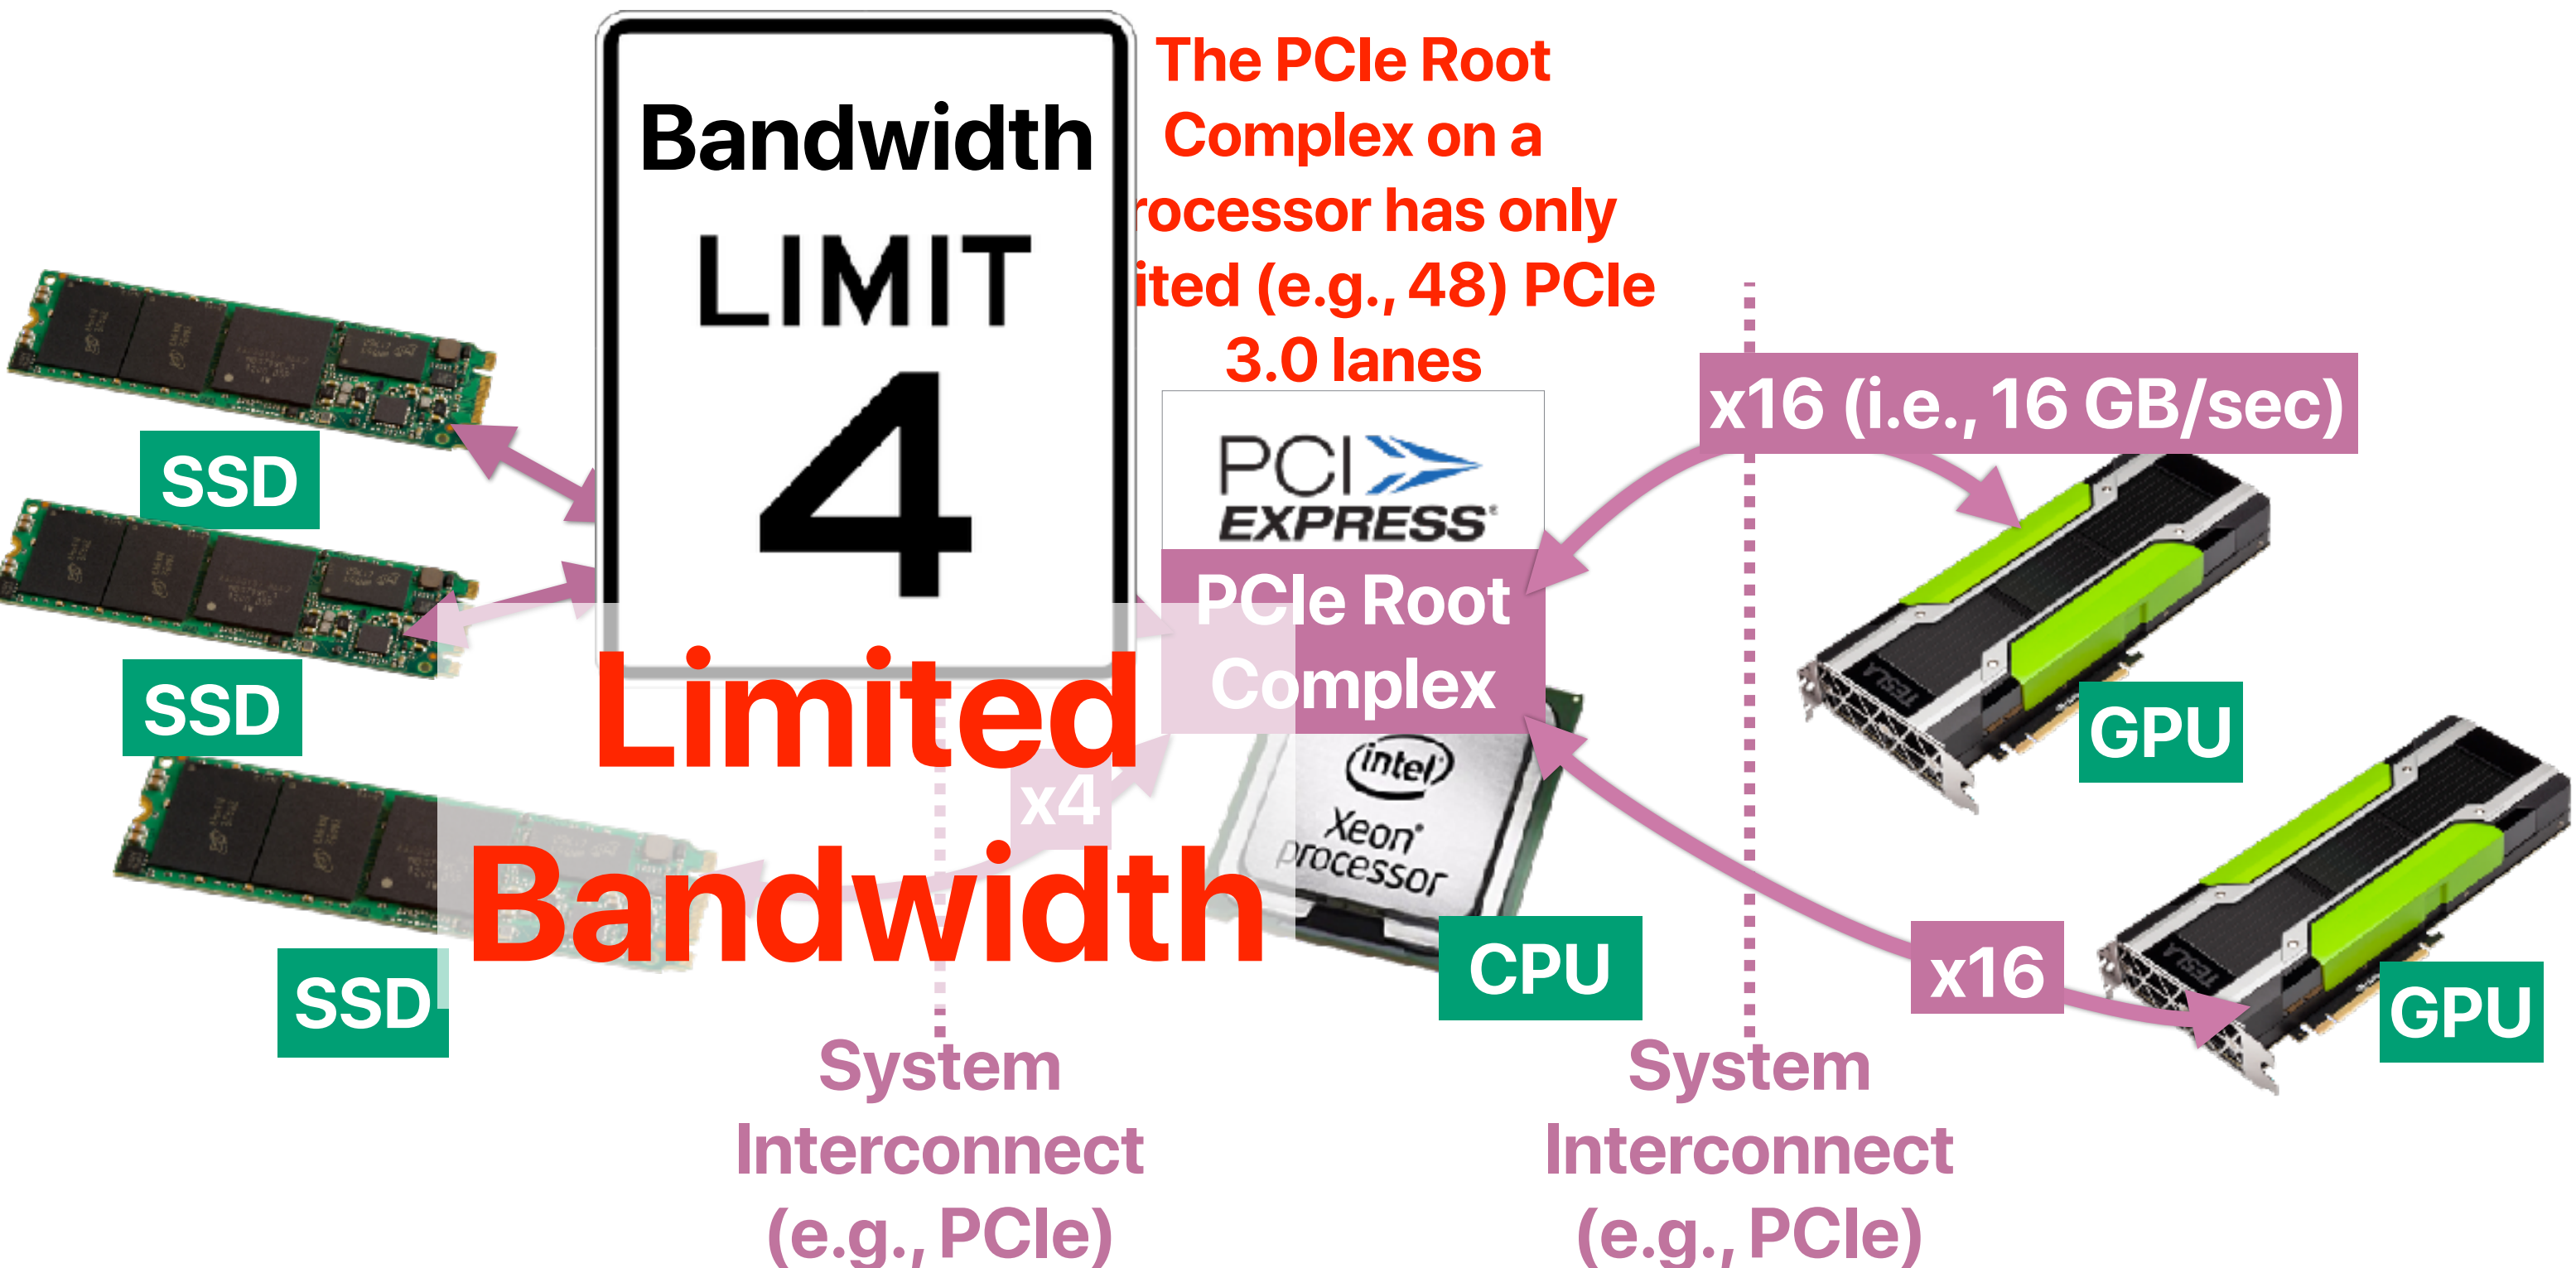
GPU-accelerated architecture



Recap: Multi-threaded I/O helps




Limited Interconnect Bandwidth



Software overhead

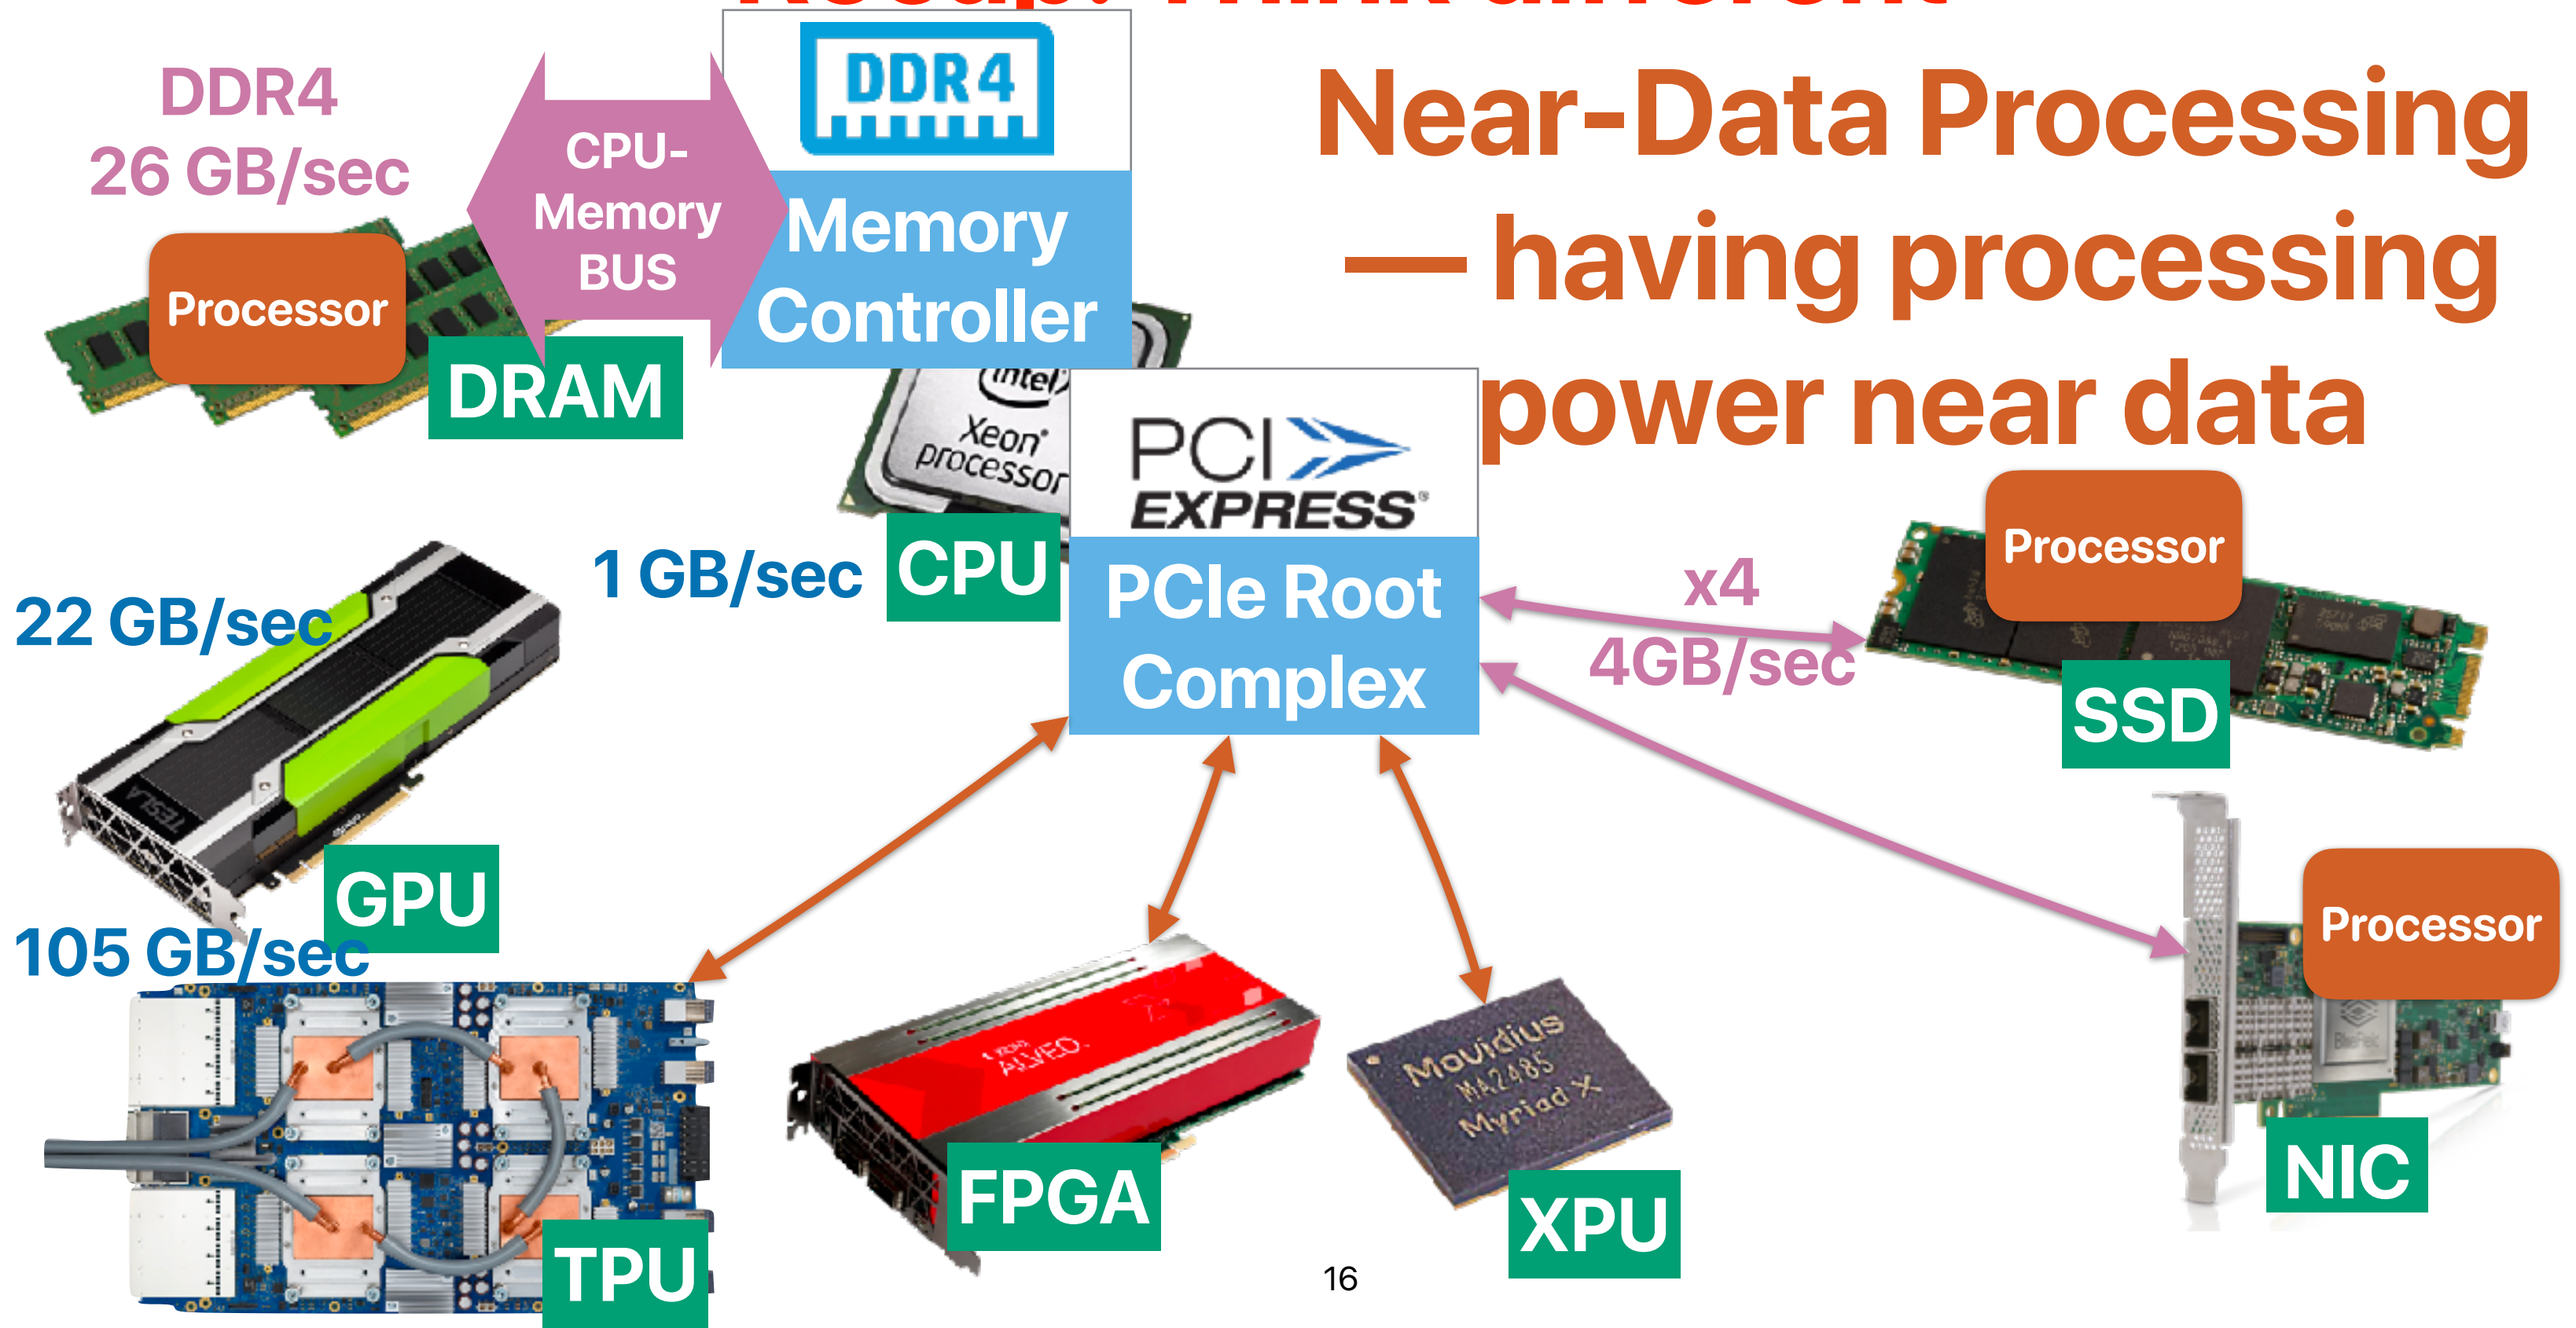
Label	Description	Baseline latency (μ s)	
		Write	Read
OS/User	OS and userspace overhead	1.98	1.95
OS/User	Linux block queue and no-op scheduler	2.51	3.74
Schedule	Get request from queue and assign tag	0.44	0.51
Copy	Data copy into DMA buffer	0.24/KB	-
Issue	PIO command writes to Moneta	1.18	1.15
DMA	DMA from host to Moneta buffer	0.93/KB	-
Ring	Data from Moneta buffer to mem ctrl	0.28/KB	-
PCM	4 KB PCM memory access	4.39	5.18
Ring	Data from mem ctrl to Moneta buffer	-	0.43/KB
DMA	DMA from Moneta buffer to host	-	0.65/KB
Wait	Thread sleep during hw	11.8	12.3
Interrupt	Driver interrupt handler	1.10	1.08
Copy	Data copy from DMA buffer	-	0.27/KB
OS/User	OS return and userspace overhead	1.98	1.95
Hardware total for 4 KB (accounting for overlap)		8.2	8.0
Software total for 4 KB (accounting for overlap)		13.3	12.2
File system additional overhead		5.8	4.2



A. M. Caulfield, A. De, J. Coburn, T. I. Mollow, R. K. Gupta and S. Swanson, "Moneta: A High-Performance Storage Array Architecture for Next-Generation, Non-volatile Memories," 2010 43rd Annual IEEE/ACM International Symposium on Microarchitecture, 2010, pp. 385-395, doi: 10.1109/MICRO.2010.33.

Recap: Think different

Near-Data Processing — having processing power near data



Where do you see the opportunity of in-storage processing in address the bandwidth limitation and software overhead?

Announcement

- What to expect in the final presentation
 - Brief recap of the why — 3-5 minutes
 - What has been tried — 5-7 minutes
 - What concrete experimental results have been measured — 5-7 minutes
 - You need to have convincing experimental setup
 - You need to have several result figures (i.e., bar charts)
 - What insights have you learned that would guide future research — 3-5 minutes

Electrical Computer Science Engineering

277

つくづく

